

```

import tempfile
from pathlib import Path

import cv2
import einops
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
import tqdm
import yaml

from helpers.adam import Adam
from modules.siren_mlp import SirenMLP

def train(config_path: Path, use_last_checkpoint: bool):
    if config_path is None:
        config_path = Path("configs/image_fit_card_f.yaml")

    # HYPERPARAMETERS
    config = yaml.safe_load(config_path.read_text())
    refresh_rate = config["display"]["refresh_rate"]
    learning_patience = config["learning"]["learning_patience"]
    learning_rates = config["learning"]["learning_rates"]
    num_iters = config["learning"]["num_iters"]
    weight_decay = config["learning"]["weight_decay"]
    num_hidden_layers = config["siren"]["num_hidden_layers"]
    hidden_layer_width = config["siren"]["hidden_layer_width"]

    rng = tf.random.get_global_generator()
    rng.reset_from_seed(0x43966E87BD57227011B5B03B58785EC1)
    tf.random.set_seed(0x43966E87BD57227011B5B03B58785EC1)

    siren = SirenMLP(
        2,
        3,
        num_hidden_layers=num_hidden_layers,
        hidden_layer_width=hidden_layer_width,
        hidden_activation=tf.math.sin,
        output_activation=tf.math.sin,
    )

    # Load the image
    input_image = cv2.imread("data/TestCardF.jpg")

    # Resize the image
    resized_img = cv2.resize(input_image, (180, 180))

    # normalize the image
    img = resized_img / 255

    target = einops.rearrange(img, "h w c -> (h w) c")

    resolution = img.shape[0]
    # Generate a linear space from -1 to 1 with the same size as the resolution
    tmp = np.linspace(-1, 1, resolution)

    # Create a meshgrid for pixel coordinates
    x, y = np.meshgrid(tmp, tmp)

    # Reshape and concatenate x and y, and cast them to float32
    x_resaped = x.reshape(-1, 1)
    y_resaped = y.reshape(-1, 1)
    img = tf.cast(tf.concat((x_resaped, y_resaped), 1), tf.float32)

    used_patience = 0

```

```

minimum_train_loss = np.inf
minimum_loss_step_num = 0

# Used For Plotting
y_train_batch_loss = np.array([])
x_train_loss_iterations = np.array([])

learning_rate_change_steps = np.array([])

# Index of the current learning rate, used to change the learning rate
# when the training loss stops improving
learning_rate_index = 0
adam = Adam(
    learning_rates[learning_rate_index],
    weight_decay=weight_decay,
)

# find the temp_dir with the prefix if it exists
# otherwise create a new one
temp_dir = None
for temp_dir in Path(tempfile.gettempdir()).iterdir():
    if temp_dir.is_dir() and temp_dir.name.startswith("siren_"):
        break

if not temp_dir.name.startswith("siren_"):
    temp_dir = tempfile.mkdtemp(prefix="siren_")

checkpoint = tf.train.Checkpoint(siren)
checkpoint_manager = tf.train.CheckpointManager(
    checkpoint,
    temp_dir,
    max_to_keep=1,
)
if use_last_checkpoint:
    print("\n\nRestoring from last checkpoint")
    checkpoint_manager.restore_or_initialize()

overall_log = tqdm.tqdm(total=0, position=1, bar_format="{desc}")
train_log = tqdm.tqdm(total=0, position=2, bar_format="{desc}")
bar = tqdm.trange(num_iters, position=3)

num_of_parameters = tf.math.add_n(
    [tf.math.reduce_prod(var.shape) for var in siren.trainable_variables]
)
print(f"\nNumber of Parameters => {num_of_parameters}")

for i in bar:
    with tf.GradientTape() as tape:
        logits = siren(img)

        current_train_loss = tf.reduce_mean((logits - target) ** 2)

    # Print initial train batch loss
    if i == 0:
        print("\n\n\n\n")
        print(f"Initial Training Loss => {current_train_loss:0.4f}")

    grads = tape.gradient(current_train_loss, siren.trainable_variables)
    adam.apply_gradients(zip(grads, siren.trainable_variables))

    current_train_loss = current_train_loss.numpy()

    if current_train_loss < minimum_train_loss:
        minimum_train_loss = current_train_loss
        minimum_loss_step_num = i

```

```

        checkpoint_manager.save()
        used_patience = i - minimum_loss_step_num

        y_train_batch_loss = np.append(y_train_batch_loss, current_train_loss)
        x_train_loss_iterations = np.append(x_train_loss_iterations, i)

        if i % refresh_rate == (refresh_rate - 1):
            learning_rates_left = len(learning_rates) - learning_rate_index
            patience_left = learning_patience - used_patience
            overall_description = (
                f"Minimum Train Loss => {minimum_train_loss:0.4f}  "
                + f"Learning Rates Left => {learning_rates_left}  "
                + f"Patience Left => {patience_left}  "
            )
            overall_log.set_description_str(overall_description)
            overall_log.refresh()

            train_description = f"Train Batch Loss => {current_train_loss:0.4f}  "
            train_log.set_description_str(train_description)
            train_log.update(refresh_rate)

            bar_description = f"Step => {i}"
            bar.set_description(bar_description)
            bar.refresh()

            # if the training loss has not improved for learning_patience
            if (
                current_train_loss > minimum_train_loss
                and i - minimum_loss_step_num > learning_patience
            ):
                if learning_rate_index == (len(learning_rates) - 1):
                    break
                learning_rate_index += 1
                adam.learning_rate = learning_rates[learning_rate_index]
                learning_rate_change_steps = np.append(learning_rate_change_steps, i)

                checkpoint_manager.restore_or_initialize()

            checkpoint_manager.restore_or_initialize()

            # delete the temporary directory
            tf.io.gfile.rmtree(temp_dir)

            checkpoint_manager = tf.train.CheckpointManager(
                checkpoint, "artifacts/siren/model", max_to_keep=1
            )
            checkpoint_manager.save()

            fig, ax = plt.subplots(3, 1)
            plt.subplots_adjust(hspace=1)

            ax[0].semilogy(x_train_loss_iterations, y_train_batch_loss, label="Training Loss")
            for learning_rate_change_step in learning_rate_change_steps:
                ax[0].axvline(
                    x=learning_rate_change_step,
                    color="black",
                    linestyle="dashed",
                    label="Learning Rate Change",
                )
            ax[0].axvline(
                x=minimum_loss_step_num, color="red", linestyle="dashed", label="Minimum Loss"
            )
            ax[0].set_xlabel("Iterations")
            ax[0].set_ylabel("Loss")

```

```

    output_image = einops.rearrange(logits.numpy(), "(h w) c -> h w c", h=180, w=180
)

# reshape output to input image shape
output_image = cv2.resize(
    output_image, (input_image.shape[1], input_image.shape[0])
)

ax[1].imshow(
    input_image,
    interpolation="nearest",
)
ax[1].set_title("Ground Truth")
ax[1].axis("off")

ax[2].imshow(
    output_image,
    interpolation="nearest",
)
ax[2].set_title("Prediction")
ax[2].axis("off")

print("\n\n\n")

print(f"Stop Iteration => {i}")

fig.suptitle("Siren – Card F:")

# if the file already exists add a number to the end of the file name
# to avoid overwriting
file_index = 0
while Path(f"artifacts/siren/siren_img_{file_index}.png").exists():
    file_index += 1
fig.savefig(f"artifacts/siren/siren_img_{file_index}.png")

# Save the config file as a yaml under the same name as the image
config_path = Path(f"artifacts/siren/siren_img_{file_index}.yaml")
config_path.write_text(yaml.dump(config))

# save the model
checkpoint_manager.save()
config_path = Path(f"artifacts/siren/model/model.yaml")
config_path.write_text(yaml.dump(config))

# def test(model_path: Path):
#     if model_path is None:
#         model_path = Path("artifacts/who_bites_who/model")

#     if not model_path.exists():
#         print("Model does not exist, run the train script first")
#         return

#     config_path = Path("artifacts/who_bites_who/model/model.yaml")

#     config = yaml.safe_load(config_path.read_text())
#     context_length = config["data"]["context_length"]
#     num_heads = config["transformer"]["num_heads"]
#     model_dim = config["transformer"]["model_dim"]
#     ffn_dim = config["transformer"]["ffn_dim"]
#     num_blocks = config["transformer"]["num_blocks"]

#     rng = tf.random.get_global_generator()
#     rng.reset_from_seed(0x43966E87BD57227011B5B03B58785EC1)
#     tf.random.set_seed(0x43966E87BD57227011B5B03B58785EC1)

```

```
# vocab_file = Path("artifacts/who_bites_who/model/vocab.txt")

# transformer_decoder = TransformerDecoder(
#     context_length,
#     num_heads,
#     model_dim,
#     ffn_dim,
#     num_blocks,
#     vocab_file=vocab_file,
# )

# checkpoint = tf.train.Checkpoint(transformer_decoder)
# checkpoint.restore(tf.train.latest_checkpoint(model_path))

# print("\n\n\nEnter 'exit' to exit")
# while 1:
#     input_text = input("\n\nEnter a sentence: ")
#     if input_text == "exit":
#         break

#     tokenized_text = transformer_decoder.predict(input_text)
#     print(f"Bite Bot: " + tokenized_text)
```

```

import tensorflow as tf

class Linear(tf.Module):
    def __init__(
        self,
        num_inputs,
        num_outputs,
        bias=True,
        zero_init=False,
        siren_init=False,
    ):
        rng = tf.random.get_global_generator()

        stddev = tf.cast(tf.math.sqrt(2 / (num_inputs + num_outputs)), tf.float32)

        self.bias = bias

        w_initial_value = rng.normal(shape=[num_inputs, num_outputs], stddev=stddev)

        if zero_init:
            w_initial_value = tf.zeros(shape=[num_inputs, num_outputs])
        elif siren_init:
            w_initial_value = rng.uniform(
                minval=-tf.math.sqrt(6 / num_inputs),
                maxval=tf.math.sqrt(6 / num_inputs),
                shape=[num_inputs, num_outputs],
            )
        self.w = tf.Variable(
            w_initial_value,
            trainable=True,
            name="Linear/w",
        )

        if self.bias:
            self.b = tf.Variable(
                tf.zeros(
                    shape=[1, num_outputs],
                ),
                trainable=True,
                name="Linear/b",
            )

        # create the logits by multiplying the inputs by the weights + the
        # optional bias
        def __call__(self, x):
            z = x @ self.w

            if self.bias:
                z += self.b

            return z

```

```

import tensorflow as tf

from modules.linear import Linear

class SirenMLP(tf.Module):
    def __init__(
        self,
        num_inputs,
        num_outputs,
        num_hidden_layers=0,
        hidden_layer_width=0,
        hidden_activation=tf.identity,
        output_activation=tf.identity,
        dropout_prob=0,
        zero_init=False,
    ):
        self.num_inputs = num_inputs
        self.num_outputs = num_outputs
        self.num_hidden_layers = num_hidden_layers
        self.hidden_layer_width = hidden_layer_width
        self.hidden_activation = hidden_activation
        self.output_activation = output_activation

        self.first_linear = Linear(num_inputs, hidden_layer_width, siren_init=True)

        self.hidden_linears = [
            Linear(hidden_layer_width, hidden_layer_width, siren_init=True)
            for _ in range(self.num_hidden_layers)
        ]
        self.final_linear = Linear(
            self.hidden_layer_width,
            self.num_outputs,
            zero_init=zero_init,
            siren_init=True,
        )

        self.dropout_prob = dropout_prob

    def __call__(self, x):
        """Applies the MLP to the input

        Args:
            x (tf.tensor): input tensor of shape [batch_size, num_inputs]

        Returns:
            tf.tensor: output tensor of shape [batch_size, num_outputs]
        """
        x = self.hidden_activation(self.first_linear(x) * 30)

        for hidden_linear in self.hidden_linears:
            x = self.hidden_activation(hidden_linear(x))

        if self.dropout_prob > 0:
            x = tf.nn.dropout(x, self.dropout_prob)

        return self.output_activation(self.final_linear(x))

```

```
#!/usr/bin/env python3

import argparse
import importlib
from pathlib import Path

import argcomplete

def main():
    parser = argparse.ArgumentParser(description="Choose an example to train:")
    parser.add_argument("runner", type=Path,
                        help="Path to the runner file")
    parser.add_argument("--config", "-c", type=Path,
                        nargs='?', help="Path to the config file")
    parser.add_argument("--restore_from_checkpoint", "-r", action="store_true",
                        help="Whether or not to use the last checkpoint")

    argcomplete.autocomplete(parser)
    args = parser.parse_args()

    runner = importlib.import_module(f"runners.{args.runner.stem}")
    runner.train(args.config, args.restore_from_checkpoint)

if __name__ == "__main__":
    main()
```